

unet-kl

November 9, 2022

0.1 prepare data

20 subjects in total, 18 will be used for training and validation, 2 will be used for independent testing raw images in nifti format, conver to png format, in order to use the existing augmentation solution.

```
[1]: import os
import numpy as np
import nibabel as nib
import matplotlib.pyplot as plt
import torch.nn as nn
from PIL import Image

[2]: from fastai.data.all import *
from fastai.vision.all import *

[3]: #nii folder has orig/label images in nifti format, orig and label would have
    ↪corresponding images in png formats
train_path = Path('./data/train/')
test_path = Path('./data/test/')

[4]: def nii2png(nii_path, png_path_orig, png_path_label):
    '''convert nii images to png files, and save original images to
    ↪png_path_orig, and labeled images to png_path_label'''
    os.system('if [ ! -d {imgdir} ]; then mkdir -p {imgdir} ; fi;'.
    ↪format(imgdir = png_path_orig))
    os.system('if [ ! -d {imgdir} ]; then mkdir -p {imgdir} ; fi;'.
    ↪format(imgdir = png_path_label))

    nifti_images = sorted(get_files(nii_path,extensions='.gz')) #all nifti
    ↪images under folder: data/train/nii

    label_images = nifti_images[::2]
    raw_images = nifti_images[1::2]

    for (raw,label) in zip(raw_images, label_images):
        sub_name = raw.stem[:10] #substr, e.g., ircad_e01_
```

```

raw_data = nib.load(raw,).get_fdata().astype(np.uint8) #intensity
↳within [0.255]
label_data = nib.load(label).get_fdata().astype(np.uint8)

assert raw_data.shape == label_data.shape

print('working on subject : {}'.format(sub_name[:-1]))
for idx in range(raw_data.shape[2]):
    orig_name = sub_name + '{:03}.png'.format(idx) #e.g., ircad_e01_000.
↳png
    label_name = sub_name + '{:03}L.png'.format(idx) #e.g.,
↳ircad_e01_000L.png
    Image.fromarray(raw_data[:, :, idx]).save(png_path_orig/orig_name)
    Image.fromarray(label_data[:, :, idx]).save(png_path_label/label_name)

```

```

[5]: # for training and validation;
nii2png(nii_path = train_path/'nii',png_path_orig = train_path/'orig',
↳png_path_label=train_path/'label')

```

```

working on subject : ircad_e01
working on subject : ircad_e02
working on subject : ircad_e03
working on subject : ircad_e04
working on subject : ircad_e05
working on subject : ircad_e06
working on subject : ircad_e07
working on subject : ircad_e08
working on subject : ircad_e09
working on subject : ircad_e10
working on subject : ircad_e11
working on subject : ircad_e12
working on subject : ircad_e13
working on subject : ircad_e14
working on subject : ircad_e15
working on subject : ircad_e16
working on subject : ircad_e17
working on subject : ircad_e18

```

```

[6]: # for independent testing;
nii2png(nii_path = test_path/'nii',png_path_orig = test_path/'orig',
↳png_path_label=test_path/'label')

```

```

working on subject : ircad_e19
working on subject : ircad_e20


```

0.2 verify images

```
[7]: def show_pair_overlay(seg, raw, alpha = 0.7):  
    f, (ax1, ax2) = plt.subplots(1, 2, sharey=True)  
    ax1.imshow(raw, cmap='gray')  
    ax2.imshow(raw, cmap='jet')  
    ax2.imshow(seg, cmap='gray', interpolation='none', alpha=alpha)  
    return f, (ax1, ax2)
```


```
[8]: #verify training dataset  
orig = Image.open('data/train/orig/ircad_e01_104.png');  
label = Image.open('data/train/label/ircad_e01_104L.png')  
show_pair_overlay(label, orig, alpha = 0.7)
```

[8]: (<Figure size 640x480 with 2 Axes>, (<AxesSubplot: >, <AxesSubplot: >))



```
[9]: #verify testing dataset  
orig = Image.open('data/test/orig/ircad_e19_041.png');  
label = Image.open('data/test/label/ircad_e19_041L.png')  
show_pair_overlay(label, orig, alpha = 0.7)
```





[9]: (<Figure size 640x480 with 2 Axes>, (<AxesSubplot: >, <AxesSubplot: >))



0.3 DataLoaders

```
[141]: path = Path('./')
dls = SegmentationDataLoaders.from_label_func(
    path, bs=4, fnames = get_image_files(path/'data/train/orig'),
    label_func = lambda o: path/'data/train/label'/f'{o.stem}L{o.suffix}',
    #liver images are put at data/label, with 'L' appended to the original name
    img_cls=PILImageBW, #only one channel, gray, not rgb;
    item_tfms=RandomResizedCrop(256, min_scale=0.3), # works on each pair of
    (orig, label)
    batch_tfms=aug_transforms(flip_vert=True) # works on each batch
)
```

```
[146]: dls.show_batch()
```



```
[35]: #check batch dims
x,y = dls.train.one_batch()
x.shape, y.shape
```

```
[35]: (torch.Size([4, 1, 256, 256]), torch.Size([4, 256, 256]))
```

0.4 unet with pretrained resnet

```
[153]: loss_func = CrossEntropyLossFlat(axis=1) #softmax over the channel dim
learn = unet_learner(dls, resnet34, n_in = 1, n_out=2, opt_func = Adam,
    ↪ loss_func = loss_func, metrics=Dice())
```

```
[16]: learn.model
```

```
[16]: DynamicUnet(
  (layers): ModuleList(
    (0): Sequential(
      (0): Conv2d(1, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
```

```

bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
    (4): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (1): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (2): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (5): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (1): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (2): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (3): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (6): Sequential(
    (0): BasicBlock(

```

```

        (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (1): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (2): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (3): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )

```



```

    )
    (4): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (5): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (7): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)

```

```

        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (2): BasicBlock(
        (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(2): ReLU()
(3): Sequential(
  (0): ConvLayer(
    (0): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (1): ReLU()
  )
  (1): ConvLayer(
    (0): Conv2d(1024, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (1): ReLU()
  )
)
(4): UnetBlock(
  (shuf): PixelShuffle_ICNR(
    (0): ConvLayer(
      (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1))
      (1): ReLU()
    )
    (1): PixelShuffle(upscale_factor=2)
  )
  (bn): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv1): ConvLayer(
    (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
  )
)

```

```

        (conv2): ConvLayer(
          (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
          (1): ReLU()
        )
      (relu): ReLU()
    )
  (5): UnetBlock(
    (shuf): PixelShuffle_ICNR(
      (0): ConvLayer(
        (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1))
        (1): ReLU()
      )
      (1): PixelShuffle(upscale_factor=2)
    )
    (bn): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv1): ConvLayer(
      (0): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): ReLU()
    )
    (conv2): ConvLayer(
      (0): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): ReLU()
    )
    (relu): ReLU()
  )
  (6): UnetBlock(
    (shuf): PixelShuffle_ICNR(
      (0): ConvLayer(
        (0): Conv2d(384, 768, kernel_size=(1, 1), stride=(1, 1))
        (1): ReLU()
      )
      (1): PixelShuffle(upscale_factor=2)
    )
    (bn): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv1): ConvLayer(
      (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): ReLU()
    )
    (conv2): ConvLayer(
      (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): ReLU()
    )
    (relu): ReLU()
  )
  (7): UnetBlock(

```

```

(shuf): PixelShuffle_ICNR(
  (0): ConvLayer(
    (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(1, 1))
    (1): ReLU()
  )
  (1): PixelShuffle(upscale_factor=2)
)
(bn): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv1): ConvLayer(
  (0): Conv2d(192, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU()
)
(conv2): ConvLayer(
  (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU()
)
(relu): ReLU()
)
(8): PixelShuffle_ICNR(
  (0): ConvLayer(
    (0): Conv2d(96, 384, kernel_size=(1, 1), stride=(1, 1))
    (1): ReLU()
  )
  (1): PixelShuffle(upscale_factor=2)
)
(9): ResizeToOrig()
(10): MergeLayer()
(11): ResBlock(
  (convpath): Sequential(
    (0): ConvLayer(
      (0): Conv2d(97, 97, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): ReLU()
    )
    (1): ConvLayer(
      (0): Conv2d(97, 97, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    )
  )
  (idpath): Sequential()
  (act): ReLU(inplace=True)
)
(12): ConvLayer(
  (0): Conv2d(97, 2, kernel_size=(1, 1), stride=(1, 1))
)
(13): fastai.layers.ToTensorBase(tensor_cls=<class
'fastai.torch_core.TensorBase'>)
)

```

)

```
[144]: learn.fine_tune(10)
```


<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```
[145]: learn.recorder.plot_loss() #seems to converge and a few more epochs would be  
↪ better.
```











```
[23]: learn.show_results()
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Target/Prediction



```
[5]: mdl_name = 'unet-resnet34-liverseg'
```

```
[148]: learn.save(mdl_name) # target: path/models/mdl_name
```

```
[148]: <fastai.learner.Learner at 0x7fdd57c98c70>
```

0.5 predict on independent testing dataset

```
[24]: def dice_coef(y_true, y_pred, eps=1e-8):  
    y_true_f = torch.flatten(y_true)  
    y_pred_f = torch.flatten(y_pred)  
    intersection = torch.sum(y_true_f * y_pred_f)  
    return ((2. * intersection + eps) / (torch.sum(y_true_f) + torch.  
↪sum(y_pred_f) + eps)).item()
```

```
[6]: path = Path('./')  
dlsT = SegmentationDataLoaders.from_label_func(  
    path, bs=4, fnames = get_image_files(path/'data/test/orig'),  
    valid_pct = 0.99, #  
    label_func = lambda o: path/'data/test/label'/f'{o.stem}L{o.suffix}',  
    img_cls=PILImageBW, #only one channel, gray, not rgb;  
    #item_tfms=Resize(256, method=ResizeMethod.Squish), # change to this resize ↵  
↪reduces the performance slightly  
    item_tfms=RandomResizedCrop(256, min_scale=0.3), # works on each pair of ↵  
↪(orig, label)  
)
```

```
[7]: loss_func = CrossEntropyLossFlat(axis=1) #softmax over the channel dim  
test = unet_learner(dlsT, resnet34, n_in = 1, n_out=2, opt_func = Adam, ↵  
↪loss_func = loss_func, metrics=Dice()) #  
test.load(mdl_name)
```

```
/home/kli/miniconda3/envs/tf/lib/python3.9/site-  
packages/torchvision/models/_utils.py:208: UserWarning: The parameter  
'pretrained' is deprecated since 0.13 and may be removed in the future, please  
use 'weights' instead.  
    warnings.warn(  
/home/kli/miniconda3/envs/tf/lib/python3.9/site-  
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a  
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed  
in the future. The current behavior is equivalent to passing  
`weights=ResNet34_Weights.IMAGENET1K_V1`. You can also use  
`weights=ResNet34_Weights.DEFAULT` to get the most up-to-date weights.  
    warnings.warn(msg)
```

```
[7]: <fastai.learner.Learner at 0x7f7b701f2190>
```

```
[8]: test.freeze()
```

```
[9]: _, test_dice = test.validate()  
     print('independent testing subjects, dice: {:.6f}'.format(test_dice))
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
independent testing subjects, dice: 0.951924
```

```
[38]: test.show_results()
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```


Target/Prediction

