

unet-kl

November 10, 2022

0.1 prepare data

20 subjects in total, 18 will be used for training and validation, 2 will be used for independent testing raw images in nifti format, conver to png format, in order to use the existing augmentation solution.

```
[86]: import os
import numpy as np
import nibabel as nib
import matplotlib.pyplot as plt
import torch.nn as nn
from PIL import Image
import torch.nn as nn
```

```
[2]: from fastai.data.all import *
from fastai.vision.all import *
```

```
[3]: #nii folder has orig/label images in nifti format, orig and label would have
↳corresponding images in png formats
train_path = Path('./data/train/')
test_path = Path('./data/test/')
```

```
[4]: def nii2png(nii_path, png_path_orig, png_path_label):
    '''convert nii images to png files, and save original images to
    ↳png_path_orig, and labeled images to png_path_label'''
    os.system('if [ ! -d {imgdir} ]; then mkdir -p {imgdir} ; fi;'.
    ↳format(imgdir = png_path_orig))
    os.system('if [ ! -d {imgdir} ]; then mkdir -p {imgdir} ; fi;'.
    ↳format(imgdir = png_path_label))

    nifti_images = sorted(get_files(nii_path,extensions='.gz')) #all nifti
    ↳images under folder: data/train/nii

    label_images = nifti_images[::2]
    raw_images = nifti_images[1::2]

    for (raw,label) in zip(raw_images, label_images):
```

```

sub_name = raw.stem[:10] #substr, e.g., ircad_e01_
raw_data = nib.load(raw,).get_fdata().astype(np.uint8) #intensity
↳ within [0.255]
label_data = nib.load(label).get_fdata().astype(np.uint8)

assert raw_data.shape == label_data.shape

print('working on subject : {}'.format(sub_name[:-1]))
for idx in range(raw_data.shape[2]):
    orig_name = sub_name + '{:03}.png'.format(idx) #e.g., ircad_e01_000.
↳ png
    label_name = sub_name + '{:03}L.png'.format(idx) #e.g.,
↳ ircad_e01_000L.png
    Image.fromarray(raw_data[:, :, idx]).save(png_path_orig/orig_name)
    Image.fromarray(label_data[:, :, idx]).save(png_path_label/label_name)

```

```

[5]: # for training and validation;
nii2png(nii_path = train_path/'nii', png_path_orig = train_path/'orig',
↳ png_path_label=train_path/'label')

```

```

working on subject : ircad_e01
working on subject : ircad_e02
working on subject : ircad_e03
working on subject : ircad_e04
working on subject : ircad_e05
working on subject : ircad_e06
working on subject : ircad_e07
working on subject : ircad_e08
working on subject : ircad_e09
working on subject : ircad_e10
working on subject : ircad_e11
working on subject : ircad_e12
working on subject : ircad_e13
working on subject : ircad_e14
working on subject : ircad_e15
working on subject : ircad_e16
working on subject : ircad_e17
working on subject : ircad_e18

```

```

[6]: # for independent testing;
nii2png(nii_path = test_path/'nii', png_path_orig = test_path/'orig',
↳ png_path_label=test_path/'label')

```

```

working on subject : ircad_e19
working on subject : ircad_e20

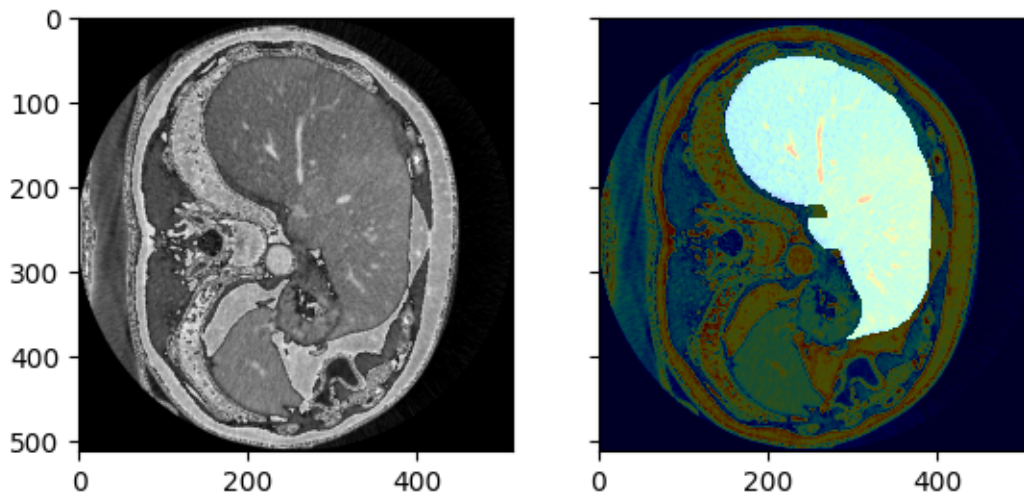
```

0.2 verify images

```
[5]: def show_pair_overlay(seg, raw, alpha = 0.7):  
      f, (ax1, ax2) = plt.subplots(1, 2, sharey=True)  
      ax1.imshow(raw, cmap='gray')  
      ax2.imshow(raw, cmap='jet')  
      ax2.imshow(seg, cmap='gray', interpolation='none', alpha=alpha)  
      return f, (ax1, ax2)
```

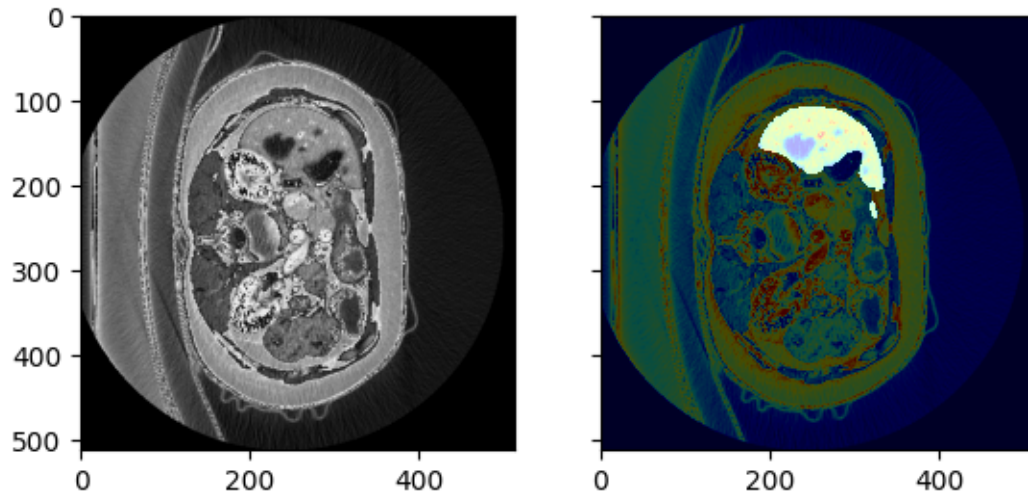
```
[6]: #verify training dataset  
orig = Image.open('data/train/orig/ircad_e01_104.png');  
label = Image.open('data/train/label/ircad_e01_104L.png')  
show_pair_overlay(label, orig, alpha = 0.7)
```

[6]: (<Figure size 640x480 with 2 Axes>, (<AxesSubplot: >, <AxesSubplot: >))



```
[7]: #verify testing dataset  
orig = Image.open('data/test/orig/ircad_e19_041.png');  
label = Image.open('data/test/label/ircad_e19_041L.png')  
show_pair_overlay(label, orig, alpha = 0.7)
```

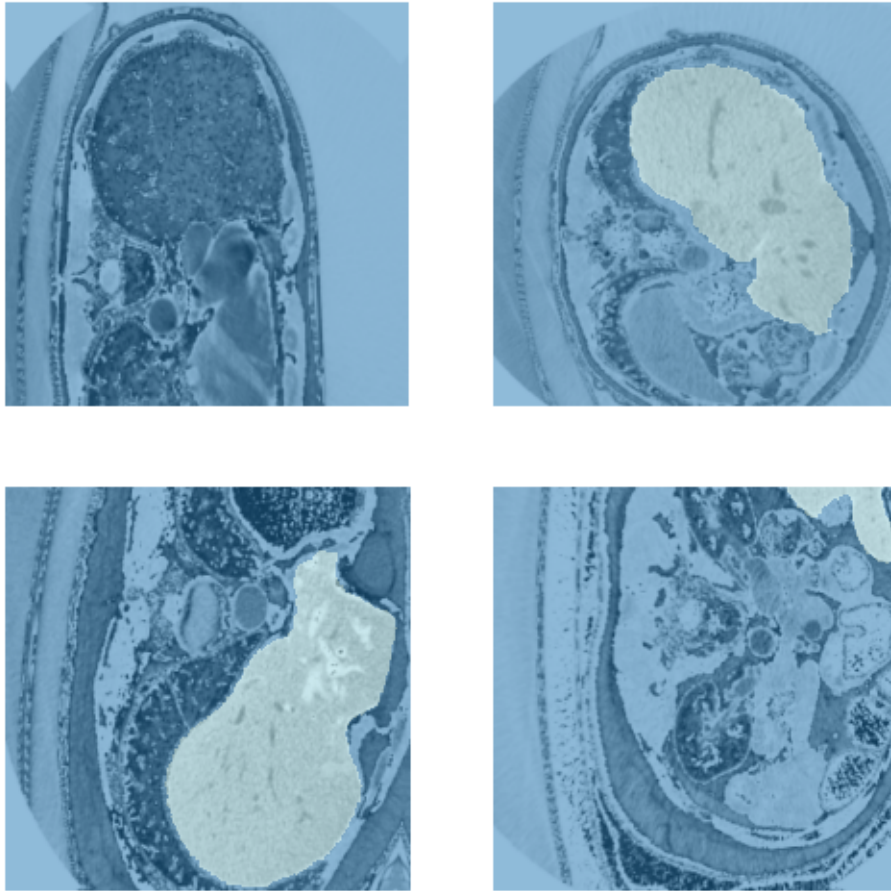
[7]: (<Figure size 640x480 with 2 Axes>, (<AxesSubplot: >, <AxesSubplot: >))



0.3 DataLoaders

```
[8]: path = Path('./')
dls = SegmentationDataLoaders.from_label_func(
    path, bs=4, fnames = get_image_files(path/'data/train/orig'),
    label_func = lambda o: path/'data/train/label'/f'{o.stem}L{o.suffix}',
    #liver images are put at data/label, with 'L' appended to the original name
    img_cls=PILImageBW, #only one channel, gray, not rgb;
    item_tfms=RandomResizedCrop(256, min_scale=0.3), # works on each pair of
    #(orig, label)
    batch_tfms=aug_transforms(flip_vert=True) # works on each batch
)
```

```
[9]: dls.show_batch()
```



```
[11]: #check batch dims
x,y = dls.train.one_batch()
x.shape, y.shape
```

```
[11]: (torch.Size([4, 1, 256, 256]), torch.Size([4, 256, 256]))
```

0.4 vanilla unet

build and test the vanilla unet with pytorch

```
[121]: def conv_block(ni, nf, kernel_size =3, padding = 1):
        return nn.Sequential(
            nn.Conv2d(ni, nf, kernel_size=3, padding=padding),
            nn.ReLU(inplace=True),
            nn.Conv2d(nf, nf, kernel_size=3, padding=padding),
            nn.ReLU(inplace=True)
        )

class VanillaUnet(nn.Module):
```

```

def __init__(self, ni=1):

    super().__init__()
    self.lconv1 = conv_block(ni, 32)
    self.lconv2 = conv_block(32, 64)
    self.lconv3 = conv_block(64, 128)
    self.lconv4 = conv_block(128, 256)
    self.lconv5 = conv_block(256, 512)

    self.rconv5 = conv_block(32, 2)
    self.rconv4 = conv_block(64+32, 32)
    self.rconv3 = conv_block(128+64, 64)
    self.rconv2 = conv_block(256+128, 128)
    self.rconv1 = conv_block(512+256, 256)

    self.maxpool = nn.MaxPool2d(2)
    self.upsample = nn.Upsample(scale_factor=2, mode='bilinear',
↪align_corners=True)

def forward(self, x):
    conv1 = self.lconv1(x)
    conv2 = self.lconv2(self.maxpool(conv1))
    conv3 = self.lconv3(self.maxpool(conv2))
    conv4 = self.lconv4(self.maxpool(conv3))
    x = self.lconv5(self.maxpool(conv4))

    x = torch.cat([self.upsample(x), conv4], dim=1)
    x = self.rconv1(x)
    x = torch.cat([self.upsample(x), conv3], dim=1)
    x = self.rconv2(x)
    x = torch.cat([self.upsample(x), conv2], dim=1)
    x = self.rconv3(x)
    x = torch.cat([self.upsample(x), conv1], dim=1)
    x = self.rconv4(x)

    out = self.rconv5(x)
    return out

```

```

[125]: loss_func=CrossEntropyLossFlat(axis=1)
learn_vun = Learner( dls, VanillaUnet(), lr=1e-3, opt_func = Adam, loss_func =
↪CrossEntropyLossFlat(axis=1), metrics=Dice(),moms=(0.95,0.85,0.95))

```

```

[127]: learn_vun.summary()

```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

[127]: VanillaUnet (Input shape: 4 x 1 x 256 x 256)

=====			
Layer (type)	Output Shape	Param #	Trainable
=====			
	4 x 32 x 256 x 256		
Conv2d		320	True
ReLU			
Conv2d		9248	True
ReLU			

	4 x 64 x 128 x 128		
Conv2d		18496	True
ReLU			
Conv2d		36928	True
ReLU			

	4 x 128 x 64 x 64		
Conv2d		73856	True
ReLU			
Conv2d		147584	True
ReLU			

	4 x 256 x 32 x 32		
Conv2d		295168	True
ReLU			
Conv2d		590080	True
ReLU			

	4 x 512 x 16 x 16		
Conv2d		1180160	True
ReLU			
Conv2d		2359808	True
ReLU			

	4 x 2 x 256 x 256		
Conv2d		578	True
ReLU			
Conv2d		38	True
ReLU			

	4 x 32 x 256 x 256		
Conv2d		27680	True
ReLU			
Conv2d		9248	True
ReLU			

```

-----
              4 x 64 x 128 x 128
Conv2d              110656      True
ReLU
Conv2d              36928      True
ReLU
-----
              4 x 128 x 64 x 64
Conv2d              442496      True
ReLU
Conv2d              147584      True
ReLU
-----
              4 x 256 x 32 x 32
Conv2d              1769728     True
ReLU
Conv2d              590080     True
ReLU
-----
              4 x 256 x 16 x 16
MaxPool2d
-----
              4 x 64 x 256 x 256
Upsample
-----

Total params: 7,846,664
Total trainable params: 7,846,664
Total non-trainable params: 0

Optimizer used: <function Adam at 0x7f071d746b80>
Loss function: FlattenedLoss of CrossEntropyLoss()

Callbacks:
- TrainEvalCallback
- CastToTensor
- Recorder
- ProgressCallback

```

```
[129]: learn_vun.fit_one_cycle(10) # the vanilla model seems to work well, though
      ↪ converged slowly
```

```
<IPython.core.display.HTML object>
```









```
<IPython.core.display.HTML object>
```

```
[132]: learn_vun.show_results()
```


<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Target/Prediction



0.5 unet with pretrained resnet

```
[8]: loss_func = CrossEntropyLossFlat(axis=1) #softmax over the channel dim
learn = unet_learner(dls, resnet34, n_in = 1, n_out=2, opt_func = Adam,
↳ loss_func = loss_func, metrics=Dice())
```

```
[16]: learn.model
```

```
[16]: DynamicUnet(
  (layers): ModuleList(
    (0): Sequential(
      (0): Conv2d(1, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
      (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): ReLU(inplace=True)
      (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
      (4): Sequential(
        (0): BasicBlock(
          (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (1): BasicBlock(
          (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (2): BasicBlock(
          (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
```

```

        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(5): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (2): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
)

```

```

    )
    (3): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (6): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (2): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)

```

```

        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (3): BasicBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (4): BasicBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (5): BasicBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (7): Sequential(
    (0): BasicBlock(
        (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),

```

```

padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (2): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
)
(1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(2): ReLU()
(3): Sequential(
  (0): ConvLayer(
    (0): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (1): ReLU()
  )
  (1): ConvLayer(
    (0): Conv2d(1024, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (1): ReLU()
  )
)

```

```

)
(4): UnetBlock(
  (shuf): PixelShuffle_ICNR(
    (0): ConvLayer(
      (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1))
      (1): ReLU()
    )
    (1): PixelShuffle(upscale_factor=2)
  )
  (bn): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv1): ConvLayer(
    (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
  )
  (conv2): ConvLayer(
    (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
  )
  (relu): ReLU()
)
(5): UnetBlock(
  (shuf): PixelShuffle_ICNR(
    (0): ConvLayer(
      (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1))
      (1): ReLU()
    )
    (1): PixelShuffle(upscale_factor=2)
  )
  (bn): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv1): ConvLayer(
    (0): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
  )
  (conv2): ConvLayer(
    (0): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
  )
  (relu): ReLU()
)
(6): UnetBlock(
  (shuf): PixelShuffle_ICNR(
    (0): ConvLayer(
      (0): Conv2d(384, 768, kernel_size=(1, 1), stride=(1, 1))
      (1): ReLU()
    )

```



```

        (1): PixelShuffle(upscale_factor=2)
    )
    (bn): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv1): ConvLayer(
        (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
    )
    (conv2): ConvLayer(
        (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
    )
    (relu): ReLU()
)
(7): UnetBlock(
    (shuf): PixelShuffle_ICNR(
        (0): ConvLayer(
            (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(1, 1))
            (1): ReLU()
        )
        (1): PixelShuffle(upscale_factor=2)
    )
    (bn): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv1): ConvLayer(
        (0): Conv2d(192, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
    )
    (conv2): ConvLayer(
        (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
    )
    (relu): ReLU()
)
(8): PixelShuffle_ICNR(
    (0): ConvLayer(
        (0): Conv2d(96, 384, kernel_size=(1, 1), stride=(1, 1))
        (1): ReLU()
    )
    (1): PixelShuffle(upscale_factor=2)
)
(9): ResizeToOrig()
(10): MergeLayer()
(11): ResBlock(
    (convpath): Sequential(
        (0): ConvLayer(
            (0): Conv2d(97, 97, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

        (1): ReLU()
    )
    (1): ConvLayer(
      (0): Conv2d(97, 97, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    )
  )
  (idpath): Sequential()
  (act): ReLU(inplace=True)
)
(12): ConvLayer(
  (0): Conv2d(97, 2, kernel_size=(1, 1), stride=(1, 1))
)
(13): fastai.layers.ToTensorBase(tensor_cls=<class
'fastai.torch_core.TensorBase'>)
)
)

```

```
[9]: learn.fine_tune(10)
```

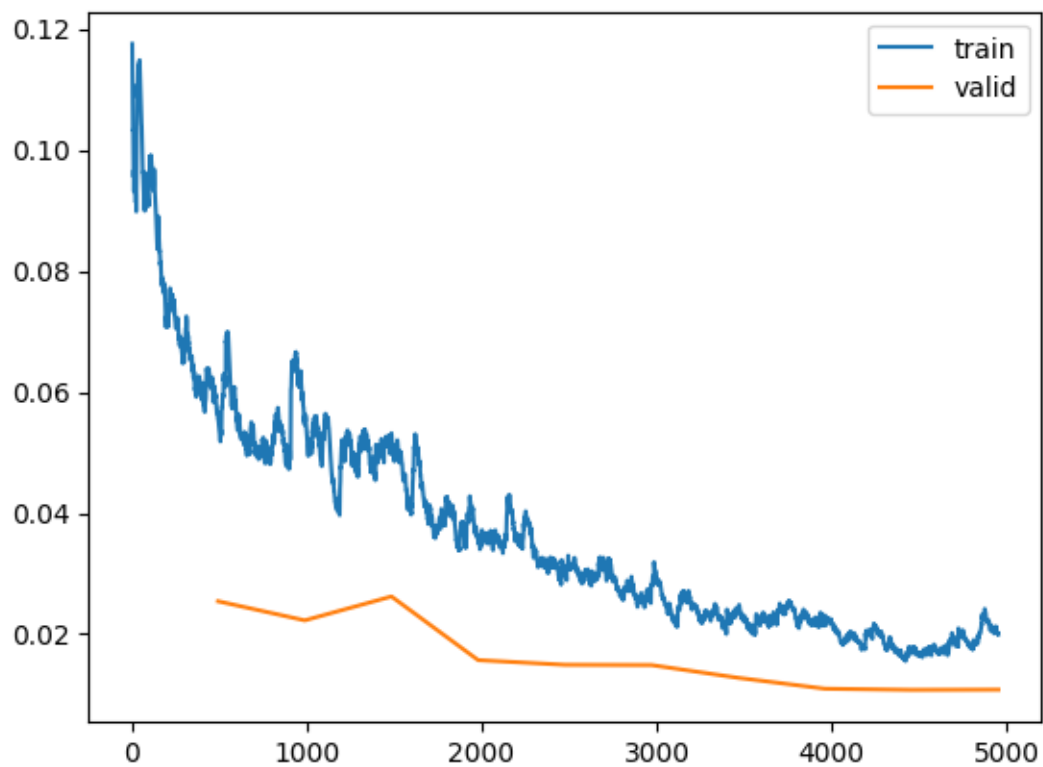
<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```
[145]: learn.recorder.plot_loss() #seems to converge and a few more epochs would be
↪ better.
```











```
[23]: learn.show_results()
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Target/Prediction



```
[12]: mdl_name = 'unet-resnet34-liverseg'
```

```
[11]: learn.save(mdl_name) # target: path/models/mdl_name
```

```
[11]: Path('models/unet-resnet34-liverseg.pth')
```

0.6 predict on independent testing dataset

```
[24]: def dice_coef(y_true, y_pred, eps=1e-8):  
    y_true_f = torch.flatten(y_true)  
    y_pred_f = torch.flatten(y_pred)  
    intersection = torch.sum(y_true_f * y_pred_f)  
    return ((2. * intersection + eps) / (torch.sum(y_true_f) + torch.  
↪sum(y_pred_f) + eps)).item()
```

```
[10]: path = Path('./')  
dlsT = SegmentationDataLoaders.from_label_func(  
    path, bs=4, fnames = get_image_files(path/'data/test/orig'),  
    valid_pct = 0.99, #  
    label_func = lambda o: path/'data/test/label'/f'{o.stem}L{o.suffix}',  
    img_cls=PILImageBW, #only one channel, gray, not rgb;  
    #item_tfms=Resize(256, method=ResizeMethod.Squish), # change to this resize  
↪reduces the performance slightly  
    item_tfms=RandomResizedCrop(256, min_scale=0.3), # works on each pair of  
↪(orig, label)  
)
```

```
[13]: loss_func = CrossEntropyLossFlat(axis=1) #softmax over the channel dim  
test = unet_learner(dlsT, resnet34, n_in = 1, n_out=2, opt_func = Adam,  
↪loss_func = loss_func, metrics=Dice()) #  
test.load(mdl_name)
```

```
[13]: <fastai.learner.Learner at 0x7f9f82dda3d0>
```

```
[8]: test.freeze()
```

```
[9]: _, test_dice = test.validate()  
print('independent testing subjects, dice: {:.6f}'.format(test_dice))
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

independent testing subjects, dice: 0.951924

```
[38]: test.show_results()
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

Target/Prediction

